

# Final Project Report-Team 6

October 23, 2023

```
[ ]: PROJECT = "aai520-project" # @param {type:"string"}
GROUP = "aai520-group6" # @param {type:"string"}
MODELS_STR = "distilbert-base-uncased, bert-base-uncased, roberta-base" #_
↳@param {type:"string"}
MODELS = [model.strip() for model in MODELS_STR.split(',')]
TASK = "question-answering" # @param {type:"string"}
DATASET = "squad_v2" # @param {type:"string"}
NUM_TRAIN_EPOCHS = 4 # @param {type:"integer"}
LEARNING_RATE = 2e-5 # @param
GRADIENT_ACCUMULATION_STEPS = 4 # @param {type:"integer"}
PER_DEVICE_TRAIN_BATCH_SIZE = 128 # @param {type:"integer"}
PER_DEVICE_EVAL_BATCH_SIZE = 128 # @param {type:"integer"}
EVALUATION_STRATEGY = 'steps' # @param {type:"string"}
EVAL_STEPS = 100 # @param {type:"integer"}
SAVE_STRATEGY = "steps" # @param {type:"string"}
SAVE_STEPS = 100 # @param {type:"integer"}
LOGGING_STEPS = 100 # @param {type:"integer"}
FP16 = True # @param {type:"boolean"}
DATALOADER_NUM_WORKERS = 2 # @param {type:"integer"}
REPORT_TO = 'tensorboard' # @param {type:"string"}
LOAD_BEST_MODEL_AT_END = True # @param {type:"boolean"}
DISABLE_TQDM = False # @param {type:"boolean"}
PUSH_TO_HUB = True # @param {type:"boolean"}
OVERWRITE_OUTPUT_DIRECTORY = True # @param {type:"boolean"}
```

```
[ ]: #-----#
# CLEAN ENVIRONMENT
#-----#

import gc
gc.collect()

import torch
torch.cuda.empty_cache()
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

#-----#
```

```

# MOUNT GOOGLE DRIVE
#-----#

from pathlib import Path
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
DRIVE_PATH = Path("/content/drive/My Drive/Colab Notebooks")

#-----#
# HANDLE PATHS
#-----#

PROJECT_PATH = DRIVE_PATH / PROJECT
PROJECT_PATH.mkdir(parents=True, exist_ok=True)
CHECKPOINTS_DIR = PROJECT_PATH / "checkpoints"
CHECKPOINTS_DIR.mkdir(parents=True, exist_ok=True)
LOGS_DIR = PROJECT_PATH / "logs"
LOGS_DIR.mkdir(parents=True, exist_ok=True)
VISUALS_DIR = PROJECT_PATH / "visuals"
VISUALS_DIR.mkdir(parents=True, exist_ok=True)

#-----#
# SYMLINK FOR FASTER FILESYSTEM TRAVERSAL
#-----#

SHORTCUT = Path("/content/project")
if not SHORTCUT.exists():
    SHORTCUT.symlink_to(PROJECT_PATH)

#-----#
# INSTALL DEPENDENCIES
#-----#

!pip install \
    -qq \
    --progress-bar=off \
    datasets \
    evaluate \
    huggingface_hub \
    python-dotenv \
    tensorboardcolab \
    transformers[torch] \
    txtai

#-----#
# IMPORT LIBRARIES & SETUP/CONFIGURE
#-----#

```

```

import json
import os

import evaluate
from datasets import load_dataset
from dotenv import load_dotenv
from evaluate.visualization import radar_plot
from huggingface_hub import EvalResult, ModelCard, ModelCardData, RepoCard,
    login
from transformers.trainer_utils import get_last_checkpoint
from txtai.pipeline import HFTrainer

# LOGIN TO HUB
load_dotenv(dotenv_path=PROJECT_PATH / ".env")
login(token=os.getenv('HF_TOKEN'))

# IGNORE TRANSFORMER WARNINGS
os.environ['TRANSFORMERS_NO_ADVISORY_WARNINGS'] = 'true'

#-----#
# LOAD DATASET
#-----#

# DATASET
ds = load_dataset(DATASET)
ds.push_to_hub(repo_id=f"{GROUP}/{DATASET}")

#-----#
# ITERATE THROUGH ALL MODELS
#-----#

combined_results = {}

for pretrained_model in MODELS:

    # RENAME MODEL TO DISTINGUISH FINETUNED MODEL
    finetuned_model = pretrained_model.replace('base', 'finetuned')
    print(f'FINE-TUNING: {finetuned_model}')

    # CONSTRUCT STANDARD IDENTIFIER FOR PUBLISHING
    IDENTIFIER = f"{GROUP}/{finetuned_model}-{DATASET}"

    # LOAD CHECKPOINTS
    model_checkpoints = CHECKPOINTS_DIR / finetuned_model

```

```

model_checkpoints.mkdir(parents=True, exist_ok=True)
last_checkpoint = get_last_checkpoint(model_checkpoints)

#-----#
# TRAIN MODEL
#-----#

# use txtai for preconfigured preprocessing of SQuAD datasets
# https://github.com/neuml/txtai
trainer = HFTrainer()
model, tokenizer = trainer(

    # SPECIFICATIONS
    base=last_checkpoint or pretrained_model,
    task=TASK,

    # DATASETS
    train=ds["train"],
    validation=ds["validation"],

    # HYPERPARAMETERS
    num_train_epochs=NUM_TRAIN_EPOCHS,
    learning_rate=LEARNING_RATE,
    gradient_accumulation_steps=GRADIENT_ACCUMULATION_STEPS,
    per_device_train_batch_size=PER_DEVICE_TRAIN_BATCH_SIZE,
    per_device_eval_batch_size=PER_DEVICE_EVAL_BATCH_SIZE,

    # CHECKPOINTS/OUTPUTS
    checkpoint=last_checkpoint,
    output_dir=str(model_checkpoints),

    # STRATEGIES
    evaluation_strategy=EVALUATION_STRATEGY,
    eval_steps=EVAL_STEPS,
    save_strategy=SAVE_STRATEGY,
    save_steps=SAVE_STEPS,
    logging_steps=LOGGING_STEPS,

    # OPTIMIZATIONS
    fp16=FP16,
    dataloader_num_workers=DATALOADER_NUM_WORKERS,

    # SAVING/LOGGING
    report_to=REPORT_TO,
    logging_dir=str(LOGS_DIR),
    load_best_model_at_end=LOAD_BEST_MODEL_AT_END,
    disable_tqdm=DISABLE_TQDM,

```

```

    # PUBLISH TO HUB
    push_to_hub=PUSH_TO_HUB,
    hub_model_id=IDENTIFIER,
    hub_token=os.getenv('HF_TOKEN')

)

#-----#
# VALIDATE MODEL
#-----#

print(f'VALIDATING: {finetuned_model}')
evaluate.logging.set_verbosity_info()
results = evaluate.evaluator("question-answering").compute(
    model_or_pipeline=IDENTIFIER,
    data=ds["validation"].select(range(2)),
    metric=DATASET,
    squad_v2_format=True,
)

for metric_type, metric_value in results.items():
    metric_name = metric_type.upper() if metric_type == 'f1' else_
↪metric_type.capitalize()
    evaluate.push_to_hub(
        model_id=IDENTIFIER,
        metric_value=metric_value,
        metric_type=metric_type,
        metric_name=metric_name,
        dataset_type=DATASET,
        dataset_name="SQuAD v2",
        dataset_split="validation",
        task_type="question-answering",
        task_name="Question Answering",
        overwrite=True
    )

# SAVE METRICS
metrics_file_path = model_checkpoints / 'metrics.json'
with open(metrics_file_path, 'w') as f:
    json.dump(results, f, indent=4)

combined_results[pretrained_model] = results

#-----#
# CREATE MODEL CARD FOR HUGGINGFACE
#-----#

```

```

# ITERATE AND ADD ALL METRICS
eval_results = [
    EvalResult(
        task_type=TASK,
        dataset_type=DATASET,
        dataset_name='SQuAD v2',
        metric_type=metric_name.replace('_', ' ').title(),
        metric_value=metric_value
    )
    for metric_name, metric_value in results.items()
]

# CREATE MODEL CARD DATA
model_card_data = ModelCardData(
    language='en',
    license='mit',
    model_name=finetuned_model,
    eval_results=eval_results,

    model_details=f"""
## Abstract
This model, '{finetuned_model}', is a question-answering chatbot
↳ trained on the SQuAD dataset, demonstrating competency in building
↳ conversational AI using recent advances in natural language processing. It
↳ utilizes a BERT model fine-tuned for extractive question answering.

## Data Collection and Preprocessing
The model was trained on the Stanford Question Answering Dataset
↳ (SQuAD), which contains over 100,000 question-answer pairs based on
↳ Wikipedia articles. The data preprocessing involved tokenizing context
↳ paragraphs and questions, truncating sequences to fit BERT's max length, and
↳ adding special tokens to mark question and paragraph segments.

## Model Architecture and Training
The architecture is based on the BERT transformer model, which was
↳ pretrained on large unlabeled text corpora. For this project, the BERT base
↳ model was fine-tuned on SQuAD for extractive question answering, with
↳ additional output layers for predicting the start and end indices of the
↳ answer span.

## SQuAD 2.0 Dataset
SQuAD 2.0 combines the existing SQuAD data with over 50,000
↳ unanswerable questions written adversarially by crowdworkers to look similar
↳ to answerable ones. This version of the dataset challenges models to not
↳ only produce answers when possible but also determine when no answer is
↳ supported by the paragraph and abstain from answering.

```

```

    """
    intended_use=f"""
    - Answering questions from the {DATASET} dataset.
    - Developing question-answering systems within the scope of the
    ↳{PROJECT}.
    - Research and experimentation in the NLP question-answering domain.
    """
    limitations_and_bias=f"""
    The model inherits limitations and biases from the '{pretrained_model}'
    ↳model, as it was trained on the same foundational data.
    It may underperform on questions that are ambiguous or too far outside
    ↳the scope of the topics covered in the {DATASET} dataset.
    Additionally, the model may reflect societal biases present in its
    ↳training data.
    """
    ethical_considerations=f"""
    This model should not be used for making critical decisions without
    ↳human oversight,
    as it can generate incorrect or biased answers, especially for topics
    ↳not covered in the training data.
    Users should also consider the ethical implications of using AI in
    ↳decision-making processes and the potential for perpetuating biases.
    """
    evaluation=f"""
    The model was evaluated on the {DATASET} dataset using various metrics.
    ↳These metrics, along with their corresponding scores,
    are detailed in the 'eval_results' section. The evaluation process
    ↳ensured a comprehensive assessment of the model's performance
    in question-answering scenarios.
    """
    training=f"""
    The model was trained over {NUM_TRAIN_EPOCHS} epochs with a learning
    ↳rate of {LEARNING_RATE}, using a batch size of {PER_DEVICE_TRAIN_BATCH_SIZE}.
    The training utilized a cross-entropy loss function and the AdamW
    ↳optimizer, with gradient accumulation over {GRADIENT_ACCUMULATION_STEPS}
    ↳steps.
    """
    tips_and_tricks=f"""
    For optimal performance, questions should be clear, concise, and
    ↳grammatically correct.
    The model performs best on questions related to topics covered in the
    ↳{DATASET} dataset.
    It is advisable to pre-process text for consistency in encoding and
    ↳punctuation, and to manage expectations for questions on topics outside the
    ↳training data.
    """

```

```

)

# CREATE AND SAVE MODEL CARD
model_card = ModelCard.from_template(model_card_data)
model_card_path = model_checkpoints / 'README.md'
model_card.save(model_card_path)

#-----#
# SAVE AND PUBLISH FINAL MODEL COMPONENTS
#-----#

# MODEL
model.save_pretrained(model_checkpoints)
model.push_to_hub(repo_id=IDENTIFIER,repo_type="model")

# TOKENIZER
tokenizer.save_pretrained(model_checkpoints)
tokenizer.push_to_hub(repo_id=IDENTIFIER, repo_type="model")

# MODEL CARD
model_card.push_to_hub(IDENTIFIER)

#-----#
# CLEAN UP BEFORE NEXT ITERATION
#-----#

del model, tokenizer, trainer
torch.cuda.empty_cache()
gc.collect()

```